

Cassy Cormier

Learning Log: Train an AI Agent to Play Flappy Bird

15698 Computer Vision 1378

Professor Anna Devarakonda

April 20th, 2025

Learning Log: Flappy Bird AI Agent Project

➤ Challenges Faced & Solutions Implemented

1. Environment Setup Failures
 - a. Early attempts to run the Flappy Bird environment in Google Colab often resulted in broken runtimes and execution errors.
 - b. Solution: Restarting the Colab runtime session consistently helped resolve these errors. Eventually, setting up the environment locally with editable package installation improved reliability.
2. Epsilon Decay & Exploration Strategy Flaws
 - a. Challenge: In Test 1, the agent's epsilon value was stuck at 1.0, meaning it kept making random moves even after 500 episodes. No learning or rewards were recorded.
 - b. Solution: The epsilon decay schedule was rechecked and implemented correctly, leading to a gradual shift from exploration to exploitation. By Test 4, epsilon had decayed to 0.1 and the agent showed improved behavior.
3. Sparse Reward Signal
 - a. Challenge: With the default reward system, the agent mostly received -5.0 for dying and almost never survived long enough to get positive feedback.
 - b. Solution: Introduced survival rewards (+0.1 per frame) to provide denser feedback and allow the agent to learn incrementally. This drastically improved learning stability.
4. Coding Bugs & Debugging Under Pressure
 - a. Challenge: Frequent bugs in the step function and inconsistent environment behavior made debugging difficult, especially with time constraints.
 - b. Solution: Leveraged ChatGPT and ClaudeAI for live debugging help, which saved time and supported deeper understanding of Python and ML code structures.
5. No Evidence of Learning
 - a. Challenge: In early tests, the agent showed no meaningful improvement in performance.
 - b. Solution: Introduced logging, performance tracking, and a reward plot function. Saved best-performing models automatically to monitor learning progress.

➤ Key Learnings & Insights

1. On Reinforcement Learning
 - a. Gained firsthand experience with Deep Q-Learning (DQN), epsilon-greedy strategies, and how crucial reward design is for agent performance.
 - b. Learned the impact of exploration-exploitation balance, reward sparsity, and target network updates in stabilizing training.
2. On Computer Vision

- a. Used image preprocessing, normalization, and grayscale conversion to reduce complexity and focus learning on relevant features like pipe edges and gaps.
 - b. Applied feature extraction and object detection techniques to enhance the agent's decision-making capability.
- 3. Personal Growth
 - a. Discovered patience and persistence are key in machine learning experiments. Debugging, especially when rewards plateau, taught resilience.
 - b. Found joy in seeing meaningful results after solving a challenging problem—like when the agent first survived multiple pipes in Test 3.
- 4. Improvements to Consider
 - a. Optimization Ideas
 - i. Incorporate hyperparameter tuning to optimize learning rate, discount factor, and batch size.
 - ii. Try Prioritized Experience Replay to speed up learning from high-value experiences.
 - b. Technical Enhancements
 - i. Explore Transfer Learning by fine-tuning pre-trained networks like MobileNetV2 for different environments.
 - ii. Add multi-agent training to simulate cooperation or competition.
- 5. Game Environment Enhancements
 - a. Introduce dynamic difficulty adjustment, like increasing pipe speed or spacing over time. Add more feedback types—like small positive rewards for flapping without dying—to reduce sparse rewards early on.
- Resources Used
 - 1. ChatGPT (OpenAI): Assisted with Python debugging and reinforcement learning concept clarification.
 - 2. ClaudeAI (Anthropic): Provided additional help with environment setup and code optimization.
 - 3. MobileNetV2 (TensorFlow/Keras): Used as a pre-trained model for feature extraction in the DQN agent.
 - 4. PyGame: Game engine for simulating the Flappy Bird environment.
 - 5. OpenCV: Used for image preprocessing (grayscale conversion, normalization, feature extraction).
 - 6. TensorFlow/Keras: Built and trained the neural network models.
 - 7. JupyterLab & Google Colab / Local Setup: Used for coding, debugging, and model training.